

## Laboratorio 2

Alvaro Frias Garay - Ary Lautaro Di Bartolo

Universidad Nacional de Córdoba - Universidad Nacional de Cuyo

2021

# Autovectorización

```
gcc -O1 -ftree-vectorize -fopt-info-vec -fopt-info-vec-missed  
wtime.c mtwister.c tiny_mc.c -lm
```

# Autovectorización

- ▶ `tiny_mc.c:107:5: missed: couldn't vectorize loop`

# Autovectorización

- ▶ `tiny_mc.c:107:5: missed: couldn't vectorize loop`
- ▶ `tiny_mc.c:107:5: missed: not vectorized: multiple nested loops.`

# Autovectorización

- ▶ `tiny_mc.c:107:5: missed: couldn't vectorize loop`
- ▶ `tiny_mc.c:107:5: missed: not vectorized: multiple nested loops.`
- ▶ `tiny_mc.c:78:16: missed: couldn't vectorize loop`

# Autovectorización

- ▶ `tiny_mc.c:107:5`: missed: couldn't vectorize loop
- ▶ `tiny_mc.c:107:5`: missed: not vectorized: multiple nested loops.
- ▶ `tiny_mc.c:78:16`: missed: couldn't vectorize loop
- ▶ `tiny_mc.c:78:16`: missed: not vectorized: control flow in loop.

# Autovectorización

- ▶ `tiny_mc.c:107:5`: missed: couldn't vectorize loop
- ▶ `tiny_mc.c:107:5`: missed: not vectorized: multiple nested loops.
- ▶ `tiny_mc.c:78:16`: missed: couldn't vectorize loop
- ▶ `tiny_mc.c:78:16`: missed: not vectorized: control flow in loop.
- ▶ `tiny_mc.c:71:9`: missed: couldn't vectorize loop

# Autovectorización

- ▶ `tiny_mc.c:107:5`: missed: couldn't vectorize loop
- ▶ `tiny_mc.c:107:5`: missed: not vectorized: multiple nested loops.
- ▶ `tiny_mc.c:78:16`: missed: couldn't vectorize loop
- ▶ `tiny_mc.c:78:16`: missed: not vectorized: control flow in loop.
- ▶ `tiny_mc.c:71:9`: missed: couldn't vectorize loop
- ▶ `tiny_mc.c:71:9`: missed: not vectorized: number of iterations cannot be computed.



# Autovectorización

```
... for (unsigned int i = 0; i < PHOTONS; ++i) {  
...     photon(r);  
... }
```

# Autovectorización

```
... for (unsigned int i = 0; i < PHOTONS; ++i) {  
...     photon(r);  
...  
...     if (weight < 0.001f) { /* roulette */  
...         if ((float) genRand(&r) > 0.1f) {  
...             break;  
...         }  
...     }  
... }
```

# Autovectorización

```
... for (unsigned int i = 0; i < PHOTONS; ++i) {  
...     photon(r);  
...  
...     if (weight < 0.001f) { /* roulette */  
...         if ((float) genRand(&r) > 0.1f) {  
...             break;  
...  
...         do {  
...             ...  
...             xi1 = 2.0f * genRand(&r) - 1.0f;  
...             xi2 = 2.0f * genRand(&r) - 1.0f;  
...             t = xi1 * xi1 + xi2 * xi2;  
...             ...  
...         } while (1.0f < t);  
...     }
```



## Vectorización Manual

# Vectorización Manual

## ► Intel Intrinsics

# Vectorización Manual

- ▶ Intel Intrinsics
- ▶ AVX / AVX-2

# Vectorización Manual

- ▶ Intel Intrinsics
- ▶ AVX / AVX-2
- ▶ Vectores de 256 bits i.e 8 floats

## Detalle de la vectorización

```
... m256 x = _mm256_set1_ps(0.0f);  
... m256 y = _mm256_set1_ps(0.0f);  
... m256 z = _mm256_set1_ps(0.0f);  
... m256 u = _mm256_set1_ps(0.0f);  
... m256 v = _mm256_set1_ps(0.0f);  
... m256 w = _mm256_set1_ps(1.0f);  
... m256 weight = _mm256_set1_ps(1.0f);
```



## Detalle de la vectorización

```
... for (;;;){
```

# Detalle de la vectorización

```
... for (;;) {
```

```
... unsigned int photon_count = 0;  
... while (photon_count < PHOTONS) {
```

## Detalle de la vectorización

```
... for (unsigned int i = 0; i < 8; ++i) {  
... | ... array_rnd[i] = (float)genRand(&r);  
... }
```

## Detalle de la vectorización

```
//float t = -logf((float)genRand(&r));
    m256 t = _mm256_set_ps(-logf(array_rnd[0]),
    ..... -logf(array_rnd[1]),
    ..... -logf(array_rnd[2]),
    ..... -logf(array_rnd[3]),
    ..... -logf(array_rnd[4]),
    ..... -logf(array_rnd[5]),
    ..... -logf(array_rnd[6]),
    ..... -logf(array_rnd[7]));
/*
x += t * u;
y += t * v;
z += t * w;
*/

// fmadd_ps(a, b, c) == (a * b) + c
x = _mm256_fmadd_ps(t, u, x);
y = _mm256_fmadd_ps(t, v, y);
z = _mm256_fmadd_ps(t, w, z);
```

# Detalle de la vectorización

tinymc sin vectorizar

```
unsigned int shell = sqrtf(x*x+y*y+z*z)*shells_per_mfp;
if (shell > SHELLS-1) {
    shell = SHELLS-1;
}
heat[shell] += (1.0f - albedo) * weight;
heat2[shell] += (1.0f - albedo) * (1.0f - albedo) * weight * weight;
weight *= albedo;
```

# Detalle de la vectorización

```
// cuadrados de números  
__m256 x_squared = _mm256_mul_ps(x, x);  
__m256 y_squared = _mm256_mul_ps(y, y);  
__m256 z_squared = _mm256_mul_ps(z, z);  
  
__m256 sum_cord = _mm256_add_ps(_mm256_add_ps(x_squared, y_squared), z_squared);
```

# Detalle de la vectorización

```
/* absorb */  
_mm256i shell_vector = _mm256_cvtps_epi32(_mm256_mul_ps(_mm256_sqrt_ps(sum_cord), shells_per_mfp));  
_mm256i max_shell_vector = _mm256_set1_epi32(SHELLS - 1);  
  
/*  
if (shell > SHELLS - 1) {  
    shell = SHELLS - 1;  
}  
*/  
shell_vector = _mm256_min_epi32(shell_vector, max_shell_vector);
```

# Detalle de la vectorización

```
/*  
heat[shell] += (1.0f - albedo) * weight;  
heat2[shell] += (1.0f - albedo) * (1.0f - albedo) * weight * weight;  
weight *= albedo;  
*/  
_mm256 helper_vector = _mm256_mul_ps(_mm256_sub_ps(ones_vector, albedo), weight);  
_mm256 helper_vector_squared = _mm256_mul_ps(helper_vector, helper_vector); /* add up squares */  
  
index_array[0] = _mm256_extract_epi32(shell_vector, 0);  
index_array[1] = _mm256_extract_epi32(shell_vector, 1);  
index_array[2] = _mm256_extract_epi32(shell_vector, 2);  
index_array[3] = _mm256_extract_epi32(shell_vector, 3);  
index_array[4] = _mm256_extract_epi32(shell_vector, 4);  
index_array[5] = _mm256_extract_epi32(shell_vector, 5);  
index_array[6] = _mm256_extract_epi32(shell_vector, 6);  
index_array[7] = _mm256_extract_epi32(shell_vector, 7);  
  
for (unsigned int i = 0; i < 8; ++i) {  
    heat[index_array[i]] += (float)helper_vector[i];  
    heat2[index_array[i]] += (float)helper_vector_squared[i];  
}
```



## Detalle de la vectorización

```
do {  
    . . . .  
    . . . xi1 = 2.0f * genRand(&r) - 1.0f;  
    . . . xi2 = 2.0f * genRand(&r) - 1.0f;  
    . . . t = xi1 * xi1 + xi2 * xi2;  
    . . . .  
} while (1.0f < t);
```

# Detalle de la vectorización

```
do {  
    for (unsigned int i = 0; i < 8; ++i) {  
        array_rnd[i] = genRand(&r);  
        array_rnd2[i] = genRand(&r);  
    }  
  
    x11 = _mm256_set_ps(2.0f * array_rnd[0] - 1.0f,  
                       2.0f * array_rnd[1] - 1.0f,  
                       2.0f * array_rnd[2] - 1.0f,  
                       2.0f * array_rnd[3] - 1.0f,  
                       2.0f * array_rnd[4] - 1.0f,  
                       2.0f * array_rnd[5] - 1.0f,  
                       2.0f * array_rnd[6] - 1.0f,  
                       2.0f * array_rnd[7] - 1.0f);  
  
    x12 = _mm256_set_ps(2.0f * array_rnd2[0] - 1.0f,  
                       2.0f * array_rnd2[1] - 1.0f,  
                       2.0f * array_rnd2[2] - 1.0f,  
                       2.0f * array_rnd2[3] - 1.0f,  
                       2.0f * array_rnd2[4] - 1.0f,  
                       2.0f * array_rnd2[5] - 1.0f,  
                       2.0f * array_rnd2[6] - 1.0f,  
                       2.0f * array_rnd2[7] - 1.0f);  
  
    // 1 == CMP_LT_OS == <  
    vec_mask = _mm256_cmp_ps(t, ones_vector, 1);  
  
    t = _mm256_blendv_ps(_mm256_add_ps(_mm256_mul_ps(x11, x11), _mm256_mul_ps(x12, x12)), t, vec_mask);  
} while (!mask_complete(vec_mask));
```

# Detalle de la vectorización

```
/*  
u = 2.0f * t - 1.0f;  
v = x11 * sqrtf((1.0f - u * u) * (1.0f / t));  
w = x12 * sqrtf((1.0f - u * u) * (1.0f / t));  
*/  
  
// mm256_fmsub_ps(a, b, c) == (a * b) - c  
m256 u = _mm256_fmsub_ps(twos_vector, t, ones_vector);  
  
m256 root = _mm256_sqrt_ps(_mm256_mul_ps(_mm256_sub_ps(ones_vector, _mm256_mul_ps(u, u)), _mm256_div_ps(ones_vector, t)));  
  
m256 v = _mm256_mul_ps(x11, root);  
  
m256 w = _mm256_mul_ps(x12, root);
```

## Detalle de la vectorización

```
if (weight < 0.001f) { /* roulette */  
    if ((float)genRand(&r) > 0.1f) {  
        break;  
    }  
    weight *= 10.0f;  
}
```


# Detalle de la vectorización

```
// weight < 0.001f  
__m256 weight_mask = _mm256_cmp_ps(weight, l_vector, _CMP_LT_OS);  
weight = _mm256_blendv_ps(weight, _mm256_mul_ps(weight, tens_vector), weight_mask);
```


# Detalle de la vectorización

```
for (unsigned int i = 0; i < 8; ++i) {  
    ... array_rnd[i] = genRand(&r);  
}  
  
    _mm256 rand_vec = _mm256_set_ps(array_rnd[0],  
    ... array_rnd[1],  
    ... array_rnd[2],  
    ... array_rnd[3],  
    ... array_rnd[4],  
    ... array_rnd[5],  
    ... array_rnd[6],  
    ... array_rnd[7]);  
  
    _mm256 roulette_mask = _mm256_cmp_ps(rand_vec, tl_vector, _CMP_GT_OS);  
  
x = _mm256_blendv_ps(x, _mm256_blendv_ps(x, zeros_vector, weight_mask), roulette_mask);  
y = _mm256_blendv_ps(y, _mm256_blendv_ps(y, zeros_vector, weight_mask), roulette_mask);  
z = _mm256_blendv_ps(z, _mm256_blendv_ps(z, zeros_vector, weight_mask), roulette_mask);  
u = _mm256_blendv_ps(u, _mm256_blendv_ps(u, zeros_vector, weight_mask), roulette_mask);  
v = _mm256_blendv_ps(v, _mm256_blendv_ps(v, zeros_vector, weight_mask), roulette_mask);  
w = _mm256_blendv_ps(w, _mm256_blendv_ps(w, ones_vector, weight_mask), roulette_mask);  
weight = _mm256_blendv_ps(weight, _mm256_blendv_ps(weight, ones_vector, weight_mask), roulette_mask);  
  
for (unsigned int i = 0; i < 8; ++i) {  
    ... if (roulette_mask[i] && weight_mask[i]) {  
    ... photon_count++;  
    ... }  
}
```

# Comparación entre vecotrización y mejor versión del lab 1



# Scaling de vectorización





# Mejoras para la vectorización

- ▶ Vectorizar el generador de números aleatorios para evitar ciclos lineales durante el programa