


Laboratorio 2

Alvaro Frias Garay - Ary Lautaro Di Bartolo

Universidad Nacional de Córdoba - Universidad Nacional de Cuyo

2021

Threads y blocks



Estrategias e implementación

- ▶ Un seed por thread
- ▶ Un fotón por thread
- ▶ heats en unified memory

Estrategias e implementación

```
__global__ void init_curand(curandState* rng_states)
{
    ... int gtid = blockDim.x * blockIdx.x + threadIdx.x;
    ... // setup a seed for every thread
    ... curand_init((unsigned long long)clock(), gtid, 0, &rng_states[gtid]);
}
```

Estrategias e implementación

```
__global__ void photon(float* global_heat, float* global_heat2, curandState* rng_states)
{
    int gtid = blockDim.x * blockIdx.x + threadIdx.x;

    // a photon per thread and if PHOTONS is not a multiple of 32 cap it
    if (gtid <= PHOTONS) {
        curandState thread_rng_state = rng_states[gtid];
```

Estrategias e implementación

```
for (;;) {  
    float t = -logf((float)curand_uniform(&thread_rng_state));  
    x += t * u;  
    y += t * v;  
    z += t * w;  
}
```

Estrategias e implementación

```
// atomic add  
atomicAdd(&global_heat[shell], (1.0f - albedo) * weight);  
atomicAdd(&global_heat2[shell], (1.0f - albedo) * (1.0f - albedo) * weight * weight);  
weight *= albedo;
```

Estrategias e implementación

```
do {  
    xi1 = 2.0f * curand_uniform(&thread_rng_state) - 1.0f;  
    xi2 = 2.0f * curand_uniform(&thread_rng_state) - 1.0f;  
    t = xi1 * xi1 + xi2 * xi2;  
} while (1.0f < t);
```


Estrategias e implementación

```
if (weight < 0.001f) {  
    if ((float)curand_uniform(&thread_rng_state) > 0.1f) {  
        break;  
    }  
}
```

Optimización al kernel

- ▶ Arrays heat de memoria `__shared__`

Optimización al kernel

```
int gtid = blockDim.x * blockIdx.x + threadIdx.x;
int tid = threadIdx.x;

__shared__ float heat_block[SHELLS];
__shared__ float heat2_blocks[SHELLS];

if (tid == 0) {
    for (unsigned int i = 0; i < SHELLS; ++i) {
        heat_block[i] = 0;
        heat2_blocks[i] = 0;
    }
}
```

Optimización al kernel

```
// atomic add  
atomicAdd(&heat_block[shell], (1.0f - albedo) * weight);  
atomicAdd(&heat2_blocks[shell], (1.0f - albedo) * (1.0f - albedo) * weight * weight);  
weight *= albedo;
```

Optimización al kernel


```
__syncthreads();

if (tid == 0) {
    for (unsigned int i = 0; i < SHELLS; ++i) {
        atomicAdd(&global_heat[i], heat_block[i]);
        atomicAdd(&global_heat2[i], heat2_blocks[i]);
    }
}
```


Optimización al kernel

```
double block_count = (PHOTONS + BLOCK_SIZE - 1) / BLOCK_SIZE;  
unsigned int total_num_threads = block_count * BLOCK_SIZE;
```

Comparación con el resto de los labs



Scaling



Posibles mejores

- ▶ Uso de arrays a nivel de warp
- ▶ Escalar el problema